

EN2550 - Fundamentals of Image Processing and Machine Vision

Assignment 01

Tharindu O.K.D.
19062R

March 4, 2022

GitHub Repository : https://github.com/dakshinatharindu/Image-Processing/blob/main/Assignment-01/190622R_a01.ipynb

Question 01

The intensity transformation is a way of mapping the intensity values of an image according to a given function. As shown in the below code, we can use the OpenCV LookUpTable (cv.LUT) method to map the intensity values.

```
1 array_1 = np.array([ i for i in range(0,51)])
2 array_2 = np.array([(155 / 100) * (i - 50) + 100 for i in range(51,151)])
3 array_3 = np.array([i for i in range(151,256)])
4 transform = np.concatenate((array_1, array_2, array_3),axis=0).astype(np.uint8)
5
6 im = cv.imread(r"emma_gray.jpg", cv.IMREAD_GRAYSCALE)
7 assert im is not None
8
9 transformed_image = cv.LUT(im, transform)
```

Listing 1: Intensity Transformation



Figure 1: Intensity Transformation

Figure 1 shows the intensity transformation function, original image, and, the transformed image. We can see that, in this transformation function, pixels with intensity values between 50 and 150, are increased but other intensities are not changed. We can see this result in the above-transformed image. The intensities of gray pixels of the original image are increased so that we can see a clear white region on her face as well as background is also changed. But dark and white pixels are not changed.

Question 02

To accentuate some particular pixels of an image, we can increase the intensity range of those pixels and reduce the intensity range of other pixels. i.e. increase the contrast of the required pixels. Then we can see more details of those pixels. Simply we can use gamma corrections to accentuate white and gray matter.

Accentuate White Matter

Here I used gamma correction with $\gamma = 4$ to accentuate white matters.

```
1 im = cv.imread(r"brain_proton_density_slice.png", cv.IMREAD_GRAYSCALE)
2 assert im is not None
3
4 white_transform = np.array([(i / 255) ** 4) * 255 for i in range(0,256)], dtype=np.uint8)
5 white_transformed_image = cv.LUT(im, white_transform)
6
```

Listing 2: Accentuate White Matter

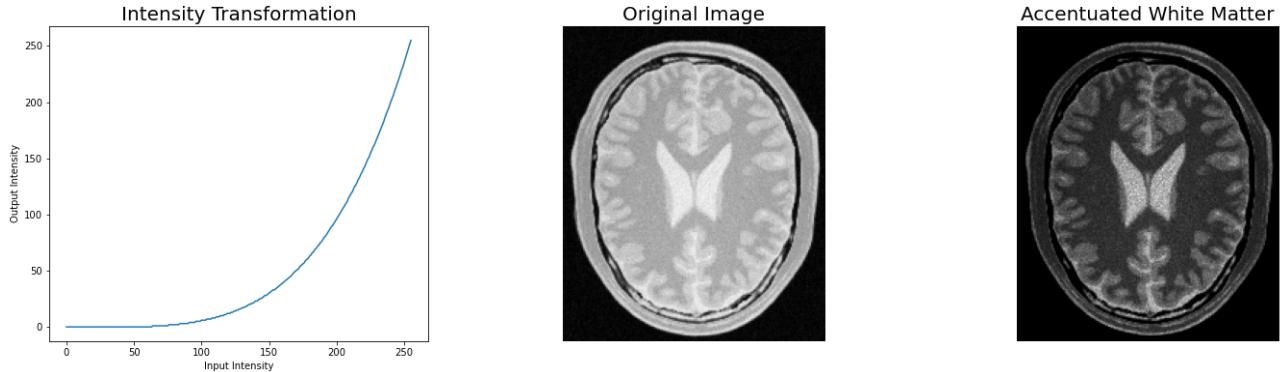


Figure 2: Accentuate White Matter

In this transformation, we can see that a wide range of dark pixels are mapped into a narrow range of dark pixels and a narrow range of white pixels are mapped into a wide range of white pixels. Therefore, we can see that the details of white matter are more highlighted compared to gray matter. Therefore, we can extract more information about white matter.

Accentuate Gray Matter

Here I used gamma correction with $\gamma = 0.3$ to accentuate gray matters.

```
1 gray_transform = np.array([(i / 255) ** 0.3) * 255 for i in range(0,256)], dtype=np.uint8)
2 gray_transformed_image = cv.LUT(im, gray_transform)
```

Listing 3: Accentuate Gray Matter

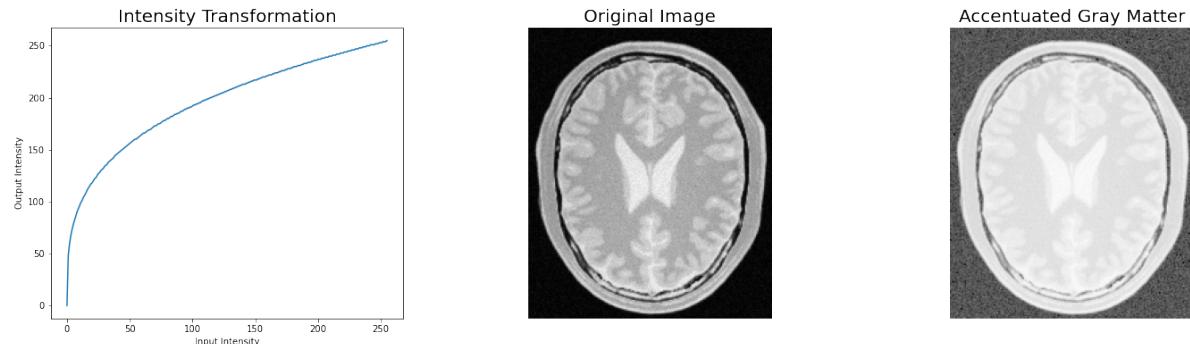


Figure 3: Accentuate Gray Matter

In this transformation, we can see that a narrow range of dark pixels are mapped into a wide range of white pixels and a wide range of white pixels are mapped into a narrow range of white pixels. Therefore, we can see that the details of gray matter are more highlighted compared to white matter. Therefore, we can extract more information about gray matter.

Question 03

In the L*a*b color space, L stands for Lightness. Here I applied a gamma correction with $\gamma = 0.4$ to the L plane.

```

1 im = cv.imread(r"highlights_and_shadows.jpg")
2 assert im is not None
3
4 LAB_image = cv.cvtColor(im, cv.COLOR_BGR2LAB)
5 original_hist = cv.calcHist([LAB_image], [0], None, [256], [0, 256])
6 gamma = 0.4
7 gamma_correction = np.array([(i / 255) ** gamma) * 255 for i in range(0, 256)], dtype=np.uint8)
8 LAB_image[:, :, 0] = cv.LUT(LAB_image[:, :, 0], gamma_correction)
9
10 corrected_hist = cv.calcHist([LAB_image], [0], None, [256], [0, 256])

```

Listing 4: Gamma Correction with $\gamma = 0.4$

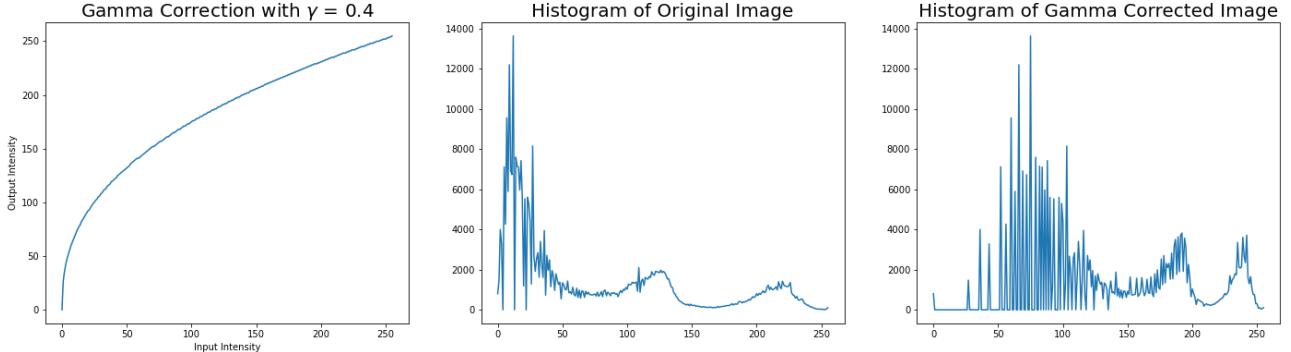


Figure 4: Histograms

When we apply gamma correction with $\gamma = 0.4$ to the L plane, the lightness of dark pixels, as well as light pixels, get increased. We can see this effect on the histogram plots. In the original image, the histogram of the L plane is more concentrated on the dark region. When applying a gamma correction to the L plane, the dark pixels are shifted to the light region. This can be seen in the Figure 4.



Figure 5: Comparison

This result can be also verified from the images in Figure 5. We can see the lightness is increased in the Gamma corrected image compared to the Original image.

Question 04

To perform the Histogram Equalization, first, we can get the histogram of the Original image using OpenCV calcHist function. Then we can get the cumulative sum of the histogram and normalize it into range (0-255). Then we can perform an intensity transformation to the original according to normalize cumulative sum to get the histogram equalized image.

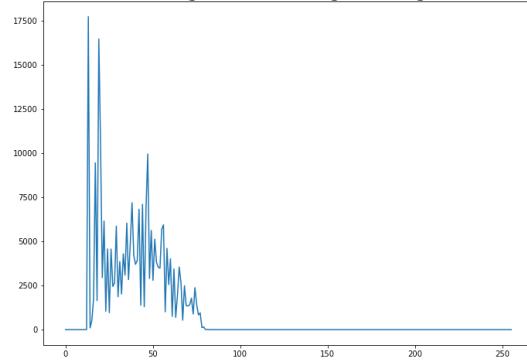
```

1 img = cv.imread(r"shells.png", cv.IMREAD_GRAYSCALE)
2 assert im is not None
3
4 original_hist = cv.calcHist([img], [0], None, [256], [0, 256])
5 cdf = original_hist.cumsum()
6 MN = img.shape[0] * img.shape[1]
7 equalize_transformation = np.array((cdf * 255) / MN, dtype=np.uint8)
8 equalize_img = cv.LUT(img, equalize_transformation)
9 equalize_hist = cv.calcHist([equalize_img], [0], None, [256], [0, 256])

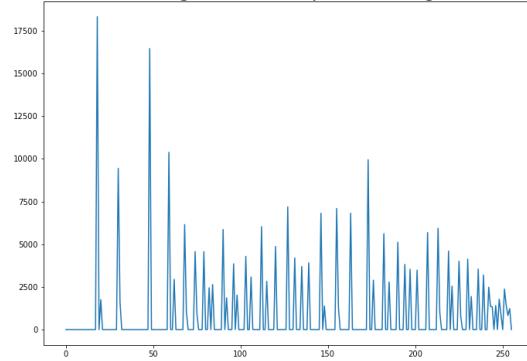
```

Listing 5: Histogram Equalization

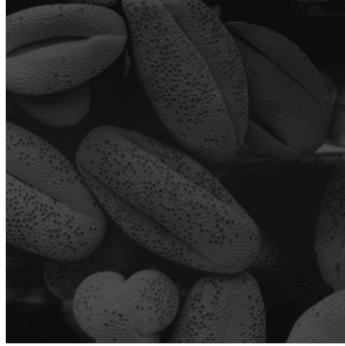
Histogram of the Original Image



Histogram of the Equalized Image



Original Image



Equalized Image

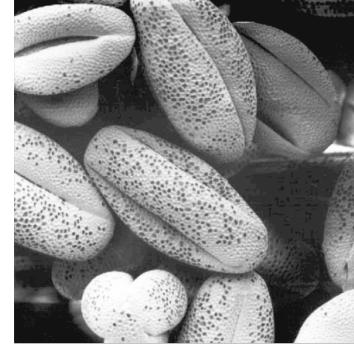


Figure 6: Histogram Equalization

We can see the histogram of the original image is more concentrated on the left side. i.e. it has more dark pixels. When applying histogram equalization, we get a more spread version of the histogram. We can see this enhancement from the above images. In the original image, the contrast is very low. We cannot get any details from it. But the equalized image has more contrast. This means we can get more details/information about the original image.

Question 05

Nearest-Neighbor Interpolation

In the Nearest-Neighbor Interpolation, we mapped each pixel of the scaled image to the nearest pixel of the original image. We can do this by iterating each pixel of the scaled image using for loops and mapping each pixel value to the nearest pixel value of the original image.

```

1 def nearestNeighborInterpolation(img, scale):
2     scaled_dimentions = (round(img.shape[0] * scale), round(img.shape[1] * scale), img.shape
3     [2])
4     scaled_img = np.zeros(scaled_dimentions, dtype=np.uint8)
5     for i in range(scaled_dimentions[0]):
6         for j in range(scaled_dimentions[1]):
7             scaled_img[i, j] = img[min(round(i / scale), img.shape[0] - 1), min(round(j /
scale), img.shape[1] - 1)]
8     return scaled_img
    
```

Listing 6: Nearest-Neighbor Interpolation

Bi-linear Interpolation

In the Bilinear Interpolation, we mapped each pixel of the scaled image to the original image by linear interpolating along the x-axis and y-axis.

```

1 def bilinearInterpolation(img, scale):
2     scaled_dimentions = (round(img.shape[0] * scale), round(img.shape[1] * scale), img.shape
3     [2])
4     scaled_img = np.zeros(scaled_dimentions, dtype=np.uint8)
5     y_ratio = (img.shape[1] - 1) / (scaled_dimentions[1] - 1)
6     x_ratio = (img.shape[0] - 1) / (scaled_dimentions[0] - 1)
7     for i in range(scaled_dimentions[0]):
8         for j in range(scaled_dimentions[1]):
9             y_floor = math.floor(j * y_ratio)
    
```

```

10     x_floor = math.floor(i * x_ratio)
11     x_ceil = min(math.ceil(i * x_ratio), img.shape[0] - 1)
12     y_weight = (j * y_ratio) - y_floor
13     x_weight = (i * x_ratio) - x_floor
14     pixel_value = img[x_floor, y_floor] * (1 - y_weight) * (1 - x_weight) + \
15         img[x_ceil, y_floor] * (1 - y_weight) * (x_weight) + \
16         img[x_floor, y_ceil] * (y_weight) * (1 - x_weight) + \
17         img[x_ceil, y_ceil] * (y_weight) * (x_weight)
18
19     scaled_img[i, j] = pixel_value
20
21 return scaled_img

```

Listing 7: Bi-linear Interpolation

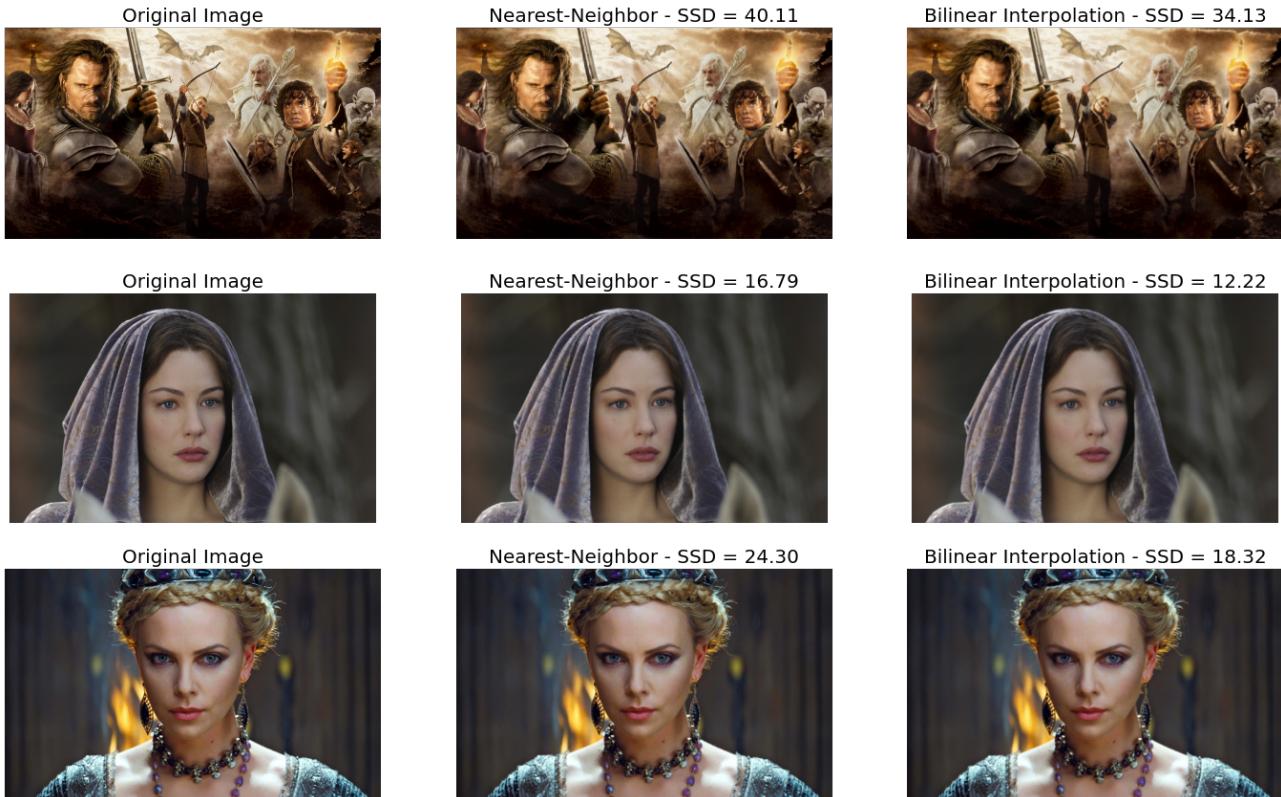


Figure 7: Comparison of Zooming

In the Nearest-Neighbor Interpolation, if we inspect closely, we can see there are small square-shaped fragments in the zoomed image. This happens because in this method we just approximate the pixel values to their nearest pixel value in the original image. Therefore we can see the same pixel is appearing in small square-shaped regions. But in the Bilinear Interpolation, we don't see this kind of behavior. This is because, in this method, we approximate the pixel value by averaging it along the x-axis and y-axis. Therefore, we can see a smooth image when applying Bilinear Interpolation. Therefore, we can say that Bilinear Interpolation performs better than Nearest-Neighbor Interpolation. This can be also verified quantitatively using the below method.

We can get a quantitative idea of the quality of these two interpolations, by calculating the Normalized Sum of Squared Difference (SSD) of two resultant images. Listing 8 shows the snipped code of the SSD function.

```

1 def SSD(img1, img2):
2     return (np.power((img1 - img2), 2)).sum() / (img1.size)

```

Listing 8: SSD

The SSD value is shown at the top of each image. We can observe that in each 3 cases, Bilinear Interpolation gives a low SSD value compared to Nearest-Neighbor Interpolation. This means Bilinear Interpolation approximates the original image better than Nearest-Neighbor Interpolation.

Question 06

Using filter2D

Sobel filtering can be used to get the edges of an image. Sobel vertical kernel extracts the horizontal edges and Sobel horizontal kernel extracts the vertical edges. In this part OpenCV filter2D function is used to apply the

kernel.

```

1 img = cv.imread(r"einstein.png", cv.IMREAD_GRAYSCALE).astype(np.float32)
2 assert im is not None
3
4 sobel_v_kernel = np.array([[-1, -2, -1], [0, 0, 0], [1, 2, 1]], dtype=np.float32)
5 f_x = cv.filter2D(img, -1, sobel_v_kernel)
6 sobel_h_kernel = np.array([[-1, 0, 1], [-2, 0, 2], [-1, 0, 1]], dtype=np.float32)
7 f_y = cv.filter2D(img, -1, sobel_h_kernel)
8 grad_mag_img = np.sqrt(f_x**2 + f_y**2)
```

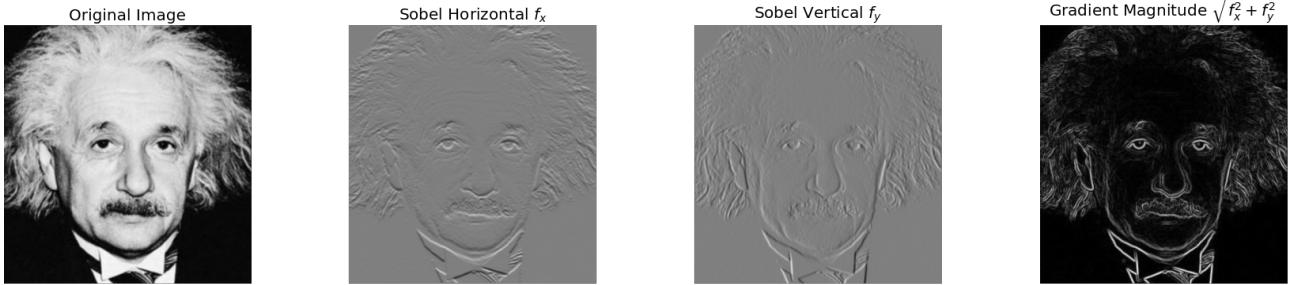


Figure 8: Using filter2D

Using Manually Coded Filter Function

This function takes the image and the kernel as the parameters. First, it rotates the kernel 180 to perform the convolution operation. Then adds the padding zeros to the boundaries of the image according to the given kernel size. After that, it applies the kernel to the image using 2 for loops.

```

1 def spacial_filter(img, kernel):
2     assert kernel.shape[0] % 2 == 1 and kernel.shape[1] % 2 == 1
3     kernel = np.rot90(np.rot90(kernel)) # rotate 180 for convolution
4     v_padding = int(kernel.shape[0] / 2)
5     h_padding = int(kernel.shape[1] / 2)
6     padded_img = np.zeros((img.shape[0] + 2 * v_padding, img.shape[1] + 2 * h_padding ), dtype =np.float32)
7     padded_img[v_padding:padded_img.shape[0] - v_padding, h_padding:padded_img.shape[1] - h_padding] = img
8     filtered_image = np.zeros(img.shape, dtype=np.float32)
9     for i in range(filtered_image.shape[0]):
10         for j in range(filtered_image.shape[1]):
11             filtered_image[i, j] = (kernel * padded_img[i: i + kernel.shape[0], j: j + kernel.shape[1]]).sum()
12     return filtered_image
13
14 f_x = spacial_filter(img, np.array([[-1, -2, -1], [0, 0, 0], [1, 2, 1]], dtype=np.float32))
15 f_y = spacial_filter(img, np.array([[-1, 0, 1], [-2, 0, 2], [-1, 0, 1]], dtype=np.float32))
16 grad_mag_img = np.sqrt(f_x**2 + f_y**2)
```

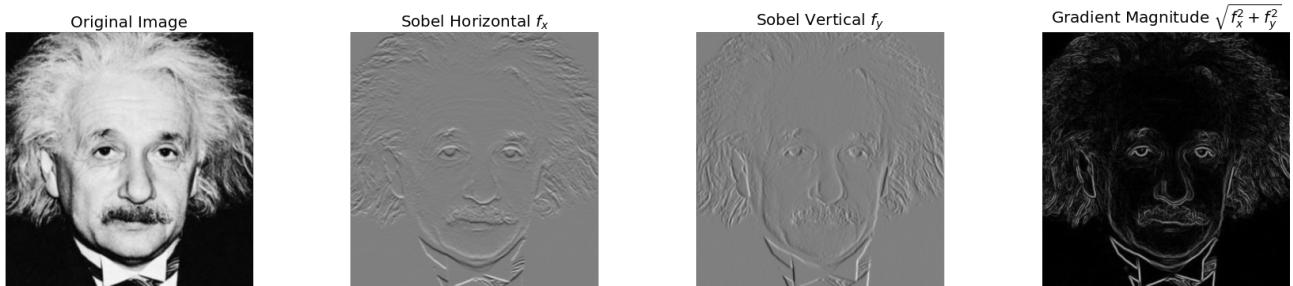


Figure 9: Using Manually Coded Filter Function

Using the Outer Product Property

We use the outer product property to perform the Sobel filtering (This can be also referred to as convolution of row vector and column vector). First, we apply the vertical kernel to the image and then apply the horizontal kernel to the resulting image. Here I used the special_filter function created by me, to apply the filter. We can also use the OpenCV filter2D function twice or the sepFilter2D function to apply the vertical kernel and horizontal kernel. This method is more efficient than the previous method because applying row kernel or column kernel needs fewer computations.

```

1 f_y = spacial_filter(img, np.array([[1], [2], [1]], dtype=np.float32))
2 f_y = spacial_filter(f_y, np.array([[1, 0, -1]], dtype=np.float32))
3
4 f_x = spacial_filter(img, np.array([[1], [0], [-1]], dtype=np.float32))
5 f_x = spacial_filter(f_x, np.array([[1, 2, 1]], dtype=np.float32))
6 grad_mag_img = np.sqrt(f_x**2 + f_y**2)

```

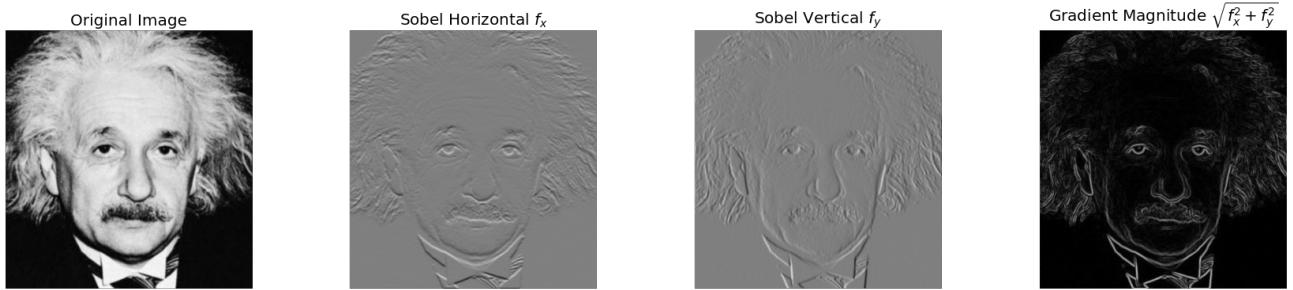


Figure 10: Using the Outer Product Property

Each 3 methods describe above gives the same results. We can see that Sobel horizontal filter extracts vertical edges of the face and Sobel vertical filter extracts horizontal edges of the face. The Gradient magnitude is calculated by squaring the Sobel vertical image and horizontal image and taking the square root. Therefore, it has vertical edges as well as horizontal edges. Therefore, we can see all the edges of the face in the gradient magnitude image.

Question 07

Segmentation

The OpenCV grabCut method can be used to segment the images. It can be initialized using 2 ways. Either using a bounding rectangle or using a mask. Since in this image we know the location of the flower, I used the bounding rectangle method.

```

1 img = cv.imread(r"daisy.jpg")
2 assert im is not None
3
4 rect = (51,147,507,441)
5 mask = np.zeros(img.shape[:2], dtype=np.uint8)
6 x,y,w,h = rect
7 mask[y:y+h, x:x+w] = 1
8
9 fgModel = np.zeros((1, 65), dtype="float")
10 bgModel = np.zeros((1, 65), dtype="float")
11 cv.grabCut(img, mask, rect, bgModel, fgModel, iterCount=10, mode=cv.GC_INIT_WITH_RECT)
12 mask2 = np.where((mask==2)|(mask==0),0,1).astype(np.uint8)
13 foreground_img = img*mask2[:, :, np.newaxis]
14 background_img = img*(1 - mask2[:, :, np.newaxis])

```

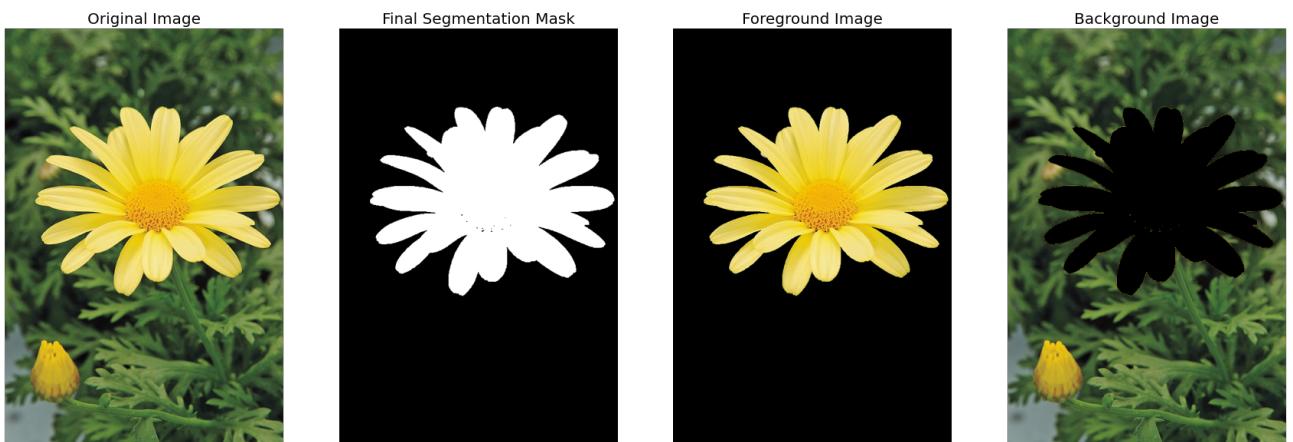


Figure 11: Segmentation

Background Blur

To blur the background I used a GaussianBlur method with kernel size = (15, 15). When we apply a filter to the background image, the cut region (black area) also gets some intensity values. If we add this blurred image directly with the foreground image we get some wrong pixel values near the edges of the flower. Therefore, before adding the background image I applied the segment mask to the blurred image to remove the unwanted pixel values at the cut region.

```
1 background_blurred_img = cv.GaussianBlur(background_img, (15, 15), 8)
2 enhanced_img = foreground_img + background_blurred_img*(1 - mask2[:, :, np.newaxis])
```

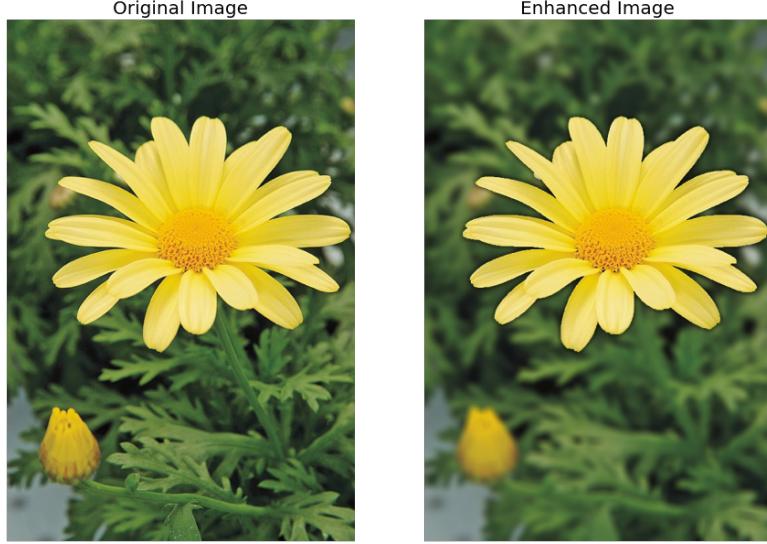


Figure 12: Background Blurred Image

Reason for Dark edges

Figure 13 shows an instant where the Gaussian kernel overlaps on an edge of segmentation. Since almost half of the kernel is in the black region, the average value of this instant is very low. i.e. the calculated intensity value will be more close to the dark region. Therefore, the pixels just beyond the edges of the flower will be darker than other pixels. When we increase the kernel size, this darker area will also increase.

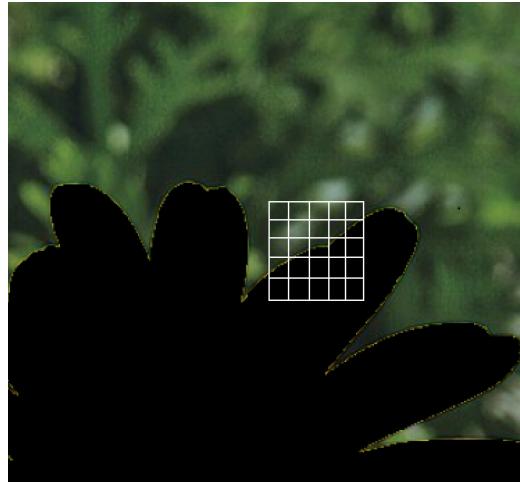


Figure 13